



CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

ETH zürich



Writing GPU Kernels

CSCS Summer School 2025

Andreas Jocksch, Prashanth Kanduri, Radim Janalik & Ben Cumming

Going Parallel : Kernels and Threads

Scaled Vector Addition (axpy)

We have used cuBLAS to perform scaled vector addition:

$$y' = \alpha x + y$$

- x and y are vectors of length n ; $x, y \in \mathbb{R}^n$
- α is scalar; $\alpha \in \mathbb{R}$

Applying `axpy` requires n operations:

$$y_i \leftarrow y_i + a * x_i, \quad i = 0, 1, \dots, n - 1$$

which can be performed **independently** and **in any order**.

Scaled Vector Addition (axpy)

We have used cuBLAS to perform scaled vector addition:

$$y' = \alpha x + y$$

- x and y are vectors of length n ; $x, y \in \mathbb{R}^n$
- α is scalar; $\alpha \in \mathbb{R}$

Applying `axpy` requires n operations:

$$y_i \leftarrow y_i + a * x_i, \quad i = 0, 1, \dots, n - 1$$

which can be performed **independently** and **in any order**.

axpy implemented on CPU with a loop

```
void axpy(double* y, const double* x, double a, int n) {
    for(int i=0; i<n; ++i) {
        y[i] = y[i] + a*x[i];
    }
}
```

Threads & Kernels

- **Threads** are streams of execution, run simultaneously on GPU.
- A **kernel** is the function run by each thread.
- CUDA provides language support for:
 - writing kernels;
 - launching many threads to execute a kernel in parallel.
- CUDA hides the low-level details of launching threads.

Threads & Kernels

- **Threads** are streams of execution, run simultaneously on GPU.
- A **kernel** is the function run by each thread.
- CUDA provides language support for:
 - writing kernels;
 - launching many threads to execute a kernel in parallel.
- CUDA hides the low-level details of launching threads.

Process for developing simple CUDA kernels

1. Formulate algorithm in terms of parallel work items.
2. Write a kernel implementing a work item on one thread.
3. Launch the kernel with the required number of threads.

Kernels

A **kernel** defines the work item for a single thread

- The work is performed by many threads executing the same kernel **simultaneously**
- Conceptually corresponds to the inner part of a loop for BLAS1 operations like `axpy`

Kernels

A **kernel** defines the work item for a single thread

- The work is performed by many threads executing the same kernel **simultaneously**
- Conceptually corresponds to the inner part of a loop for BLAS1 operations like `axpy`

host : add two vectors	CUDA : add two vectors
<pre>void add_cpu(int* a, int* b, int n){ for(auto i=0; i<n; ++i) a[i] = a[i] + b[i]; }</pre>	<pre>__global__ void add_gpu(int* a, int* b, int n){ auto i = threadIdx.x; a[i] = a[i] + b[i]; }</pre>

NOTE

- `__global__` keyword indicates a kernel
- `threadIdx` is used to find unique ID of each thread

Launching a Kernel

- Host code launches a kernel on the GPU **asynchronously**
- CUDA provides the “triple chevron” `<<<_,_>>>` syntax for launching a kernel

CPU : add two vectors	CUDA : add two vectors
<pre>auto n = 1024; auto a = host_malloc <int >(n); auto b = host_malloc <int >(n); add_cpu(a, b, n);</pre>	<pre>auto n = 1024; auto a = device_malloc <int >(n); auto b = device_malloc <int >(n); add_gpu<<<1,n>>>(a, b, n);</pre>

NOTE

`add_gpu<<<1, num_threads>>>(args...)` launches the kernel `add_gpu` with `num_threads` parallel threads

Exercise: My First Kernel

Open `axpy/axpy.cu`

1. Write a kernel that implements `axpy` for `double`
 - `axpy_kernel(double* y, double* x, double a, int n)`
 - **extra:** can you write a C++ templated version for any type?
2. Launch the kernel (look for `TODO`)
3. Compile the test and run
 - it will pass with no errors on success
 - first try with small vectors of size 8
 - try increasing launch size... what happens?
4. **extra:** can you extend the kernel to work for larger arrays?

Scaling Up : Thread Blocks

Key Observation

In the axpy exercises we were limited to 1024 threads for a kernel launch

- BUT we need to scale beyond 1024 threads for the **massive parallelism** we were promised!

Thread Blocks and Grids

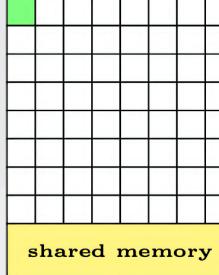
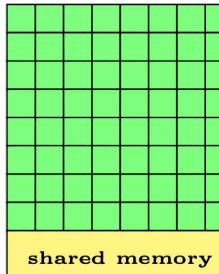
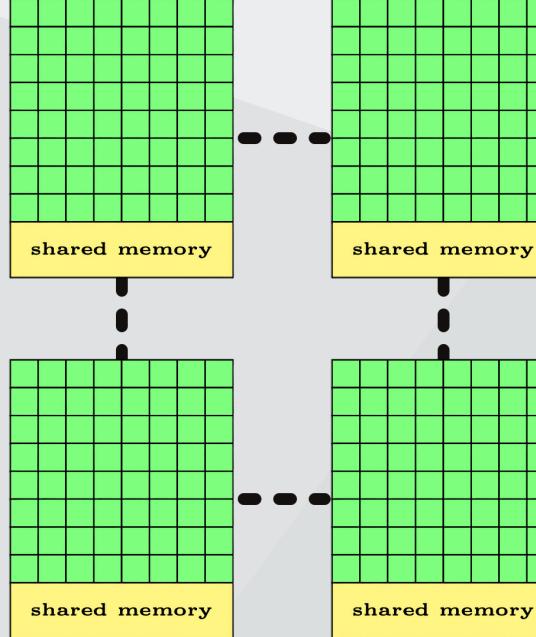
Kernels are executed in groups of threads called **thread blocks**

- the launch configuration `axpy<<<grid_dim, block_dim>>>(...)`
 - launch a **grid** of `grid_dim` **blocks**
 - each **block** has `block_dim` **threads**
 - for a total of `grid_dim × block_dim` threads
- previously we launched just one thread block `axpy<<<1, n>>>`

Why the additional complexity?

Coordination between threads doesn't scale:

- Threads in a block can synchronize and share resources
- This does not scale past a certain number of cores/threads
- Threads in a block run on the same SMX, with shared resources and thread cooperation
- Work is broken into blocks, which are distributed over the SMXs on the GPU

Concept	Hardware	Remarks
thread		<ul style="list-style-type: none"> each thread executed on one core
block		<ul style="list-style-type: none"> block executed on 1 SMX multiple blocks per SMX if sufficient resources threads in a block share SMX resources
grid		<ul style="list-style-type: none"> kernel is executed in grid of blocks blocks distributed over SMXs multiple kernels can run at same time

Calculating Thread Indices

A kernel has to calculate the index of its work item

- In `axpy` we used `threadIdx.x` for the index
- With multiple blocks, we need more information, which is available in the following **magic variables**:

Variable	Purpose
<code>gridDim</code>	total number of blocks in the grid
<code>blockDim</code>	number of threads in a thread block
<code>blockIdx</code>	index of block in the range <code>[0, gridDim-1]</code>
<code>threadIdx</code>	index of thread in thread block <code>[0, blockDim-1]</code>

Calculating Thread Indices

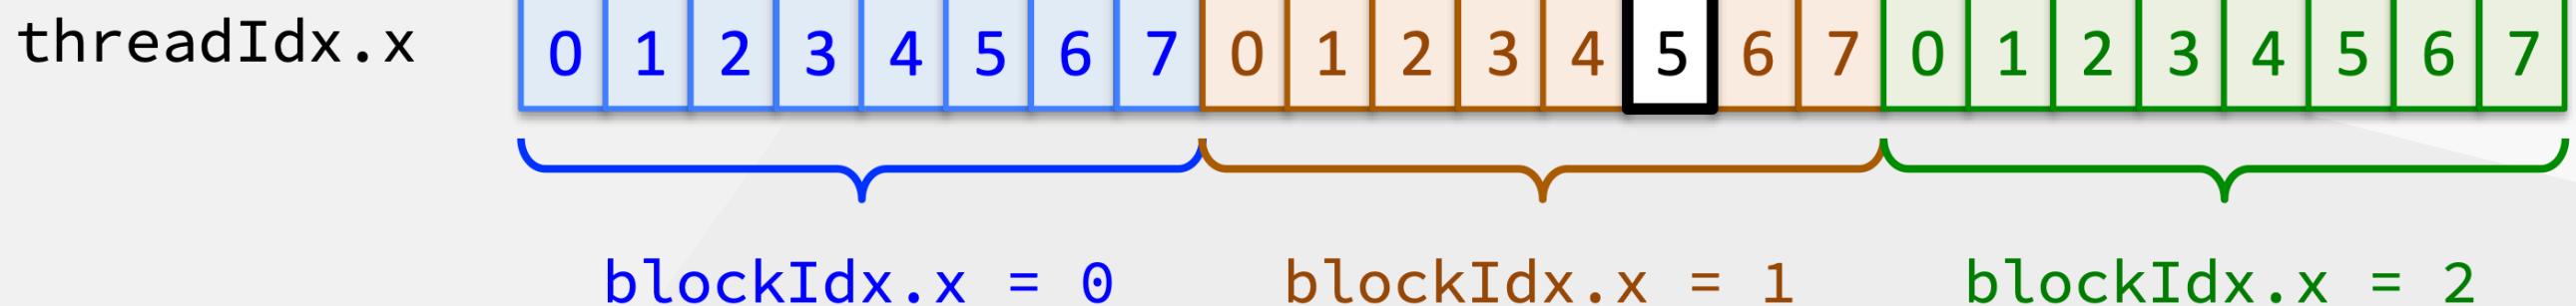
Consider accessing an array of length 24 with 8 threads per block. The **dimensions** of the kernel launch are:

- `blockDim.x == 8` (8 threads/block)
- `gridDim.x == 3` (3 blocks)

We calculate the index for our thread using the formula

```
auto index = threadIdx.x + blockDim.x*blockIdx.x
```

$$\begin{aligned} \text{index} &= \text{threadIdx.x} + \text{blockDim.x} * \text{blockIdx.x} \\ &= 5 + 8 * 1 \\ &= 13 \end{aligned}$$



Calculating Grid Dimensions

The number of thread blocks and the number of threads per block are parameters for the kernel launch:

```
kernel<<<blocks, threads_per_block>>>(...)
```

Remember to guard against overflow when the number of work items is not divisible by the thread block size.

Vector Addition with Multiple Blocks

```
__global__
void add_gpu(int* a, int* b, int n) {
    auto i = threadIdx.x + blockIdx.x*blockDim.x;
    if(i<n) { // guard against access off end of arrays
        a[i] += b[i];
    }
}
// in main ()
auto block_size = 512;
auto num_blocks = (n + (block_size-1)) / block_size;
add_gpu<<<num_blocks, block_size>>>(a, b, n);
```

Calculating Grid Dimensions

Pay attention when calculating the number of blocks in the grid, i.e. `blocks : kernel<<<blocks, threads_per_block>>>(...)`

Most likely, the number of work items `n` is not a multiple of `threads_per_block`

- some threads in the last thread block will be **idle**.

```
// in main ()  
auto block_size = 512;  
auto num_blocks = (n + block_size-1) / block_size;  
add_gpu<<<num_blocks, block_size>>>(a, b, n);
```

How many threads per block?

The number of threads per block has an impact on performance

- The optimal number depends on resources required by the kernel
- This includes registers, shared memory, computational intensity, etc

The short answer is 32, 64, or 128

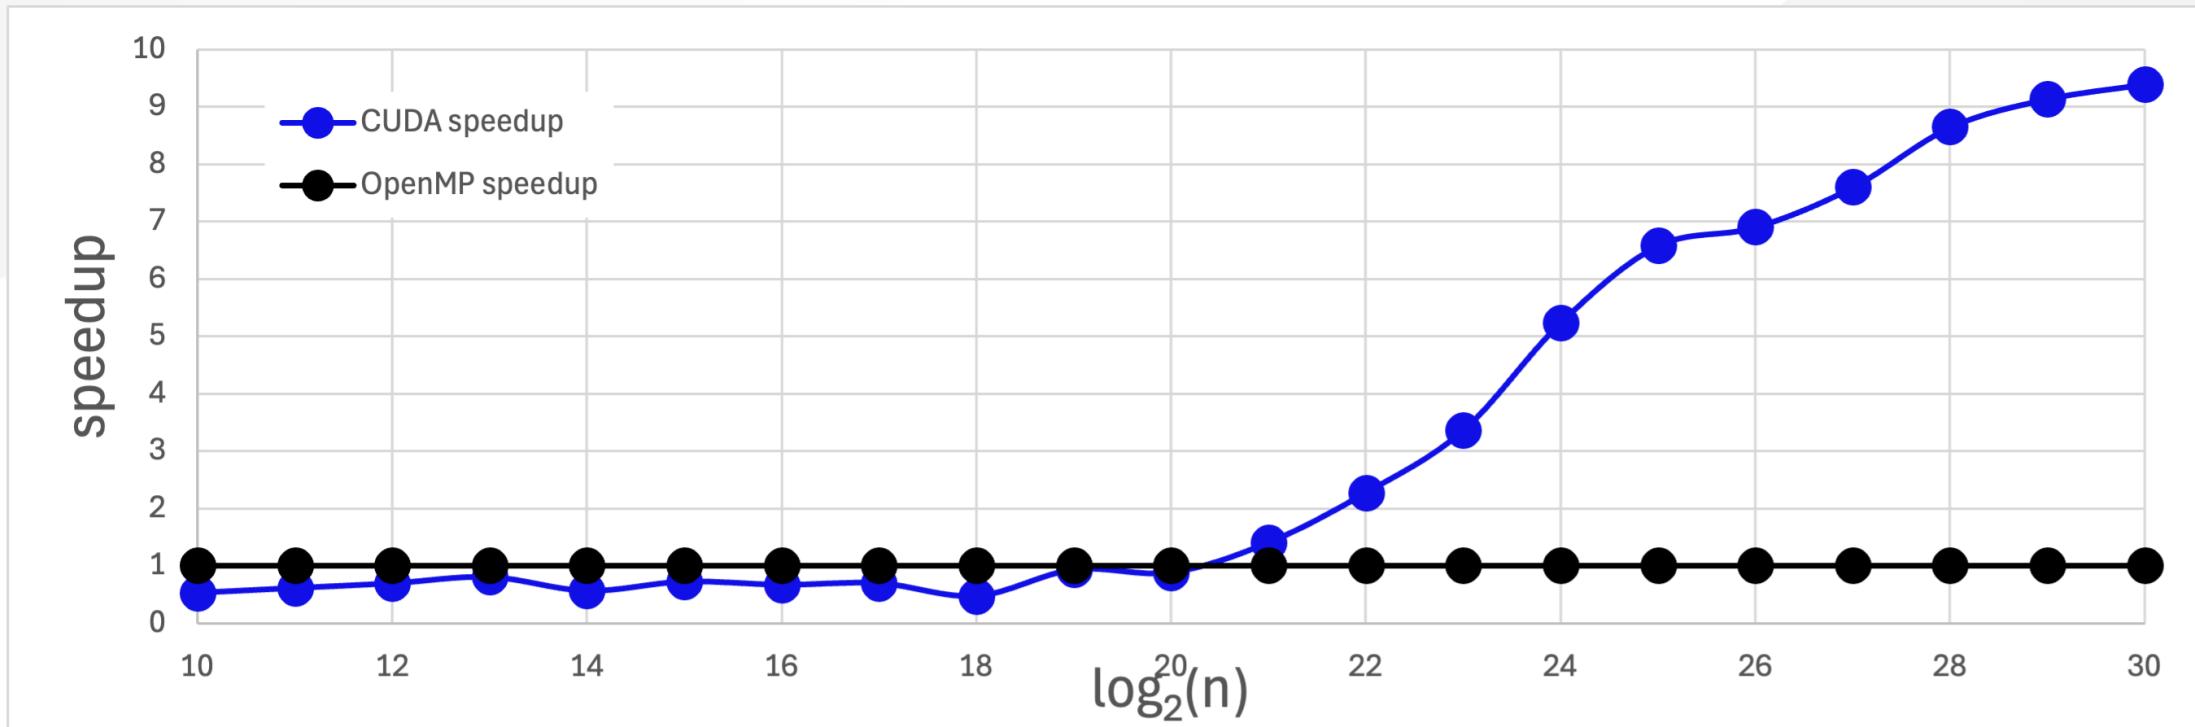
- For the main kernels in your application, perform experiments to find the ideal block size

Exercise: Blocks

Open `axpy/axpy.cu` from the last exercise

1. Extend the axpy kernel for arbitrarily large input arrays (any `n`)
2. Update the call site to calculate the grid configuration
3. Compile the test and run
 - it will pass with no errors on success
4. Experiment with varying the size of the arrays (scaling)
 - start small and increase
5. finish the `newton.cu` example
 - how do the h2d, d2h and kernel timings compare?
6. **extra:** Compare scaling with the `axpy_omp` benchmark
7. **extra:** Experiment with varying the block size

Exercise: Results



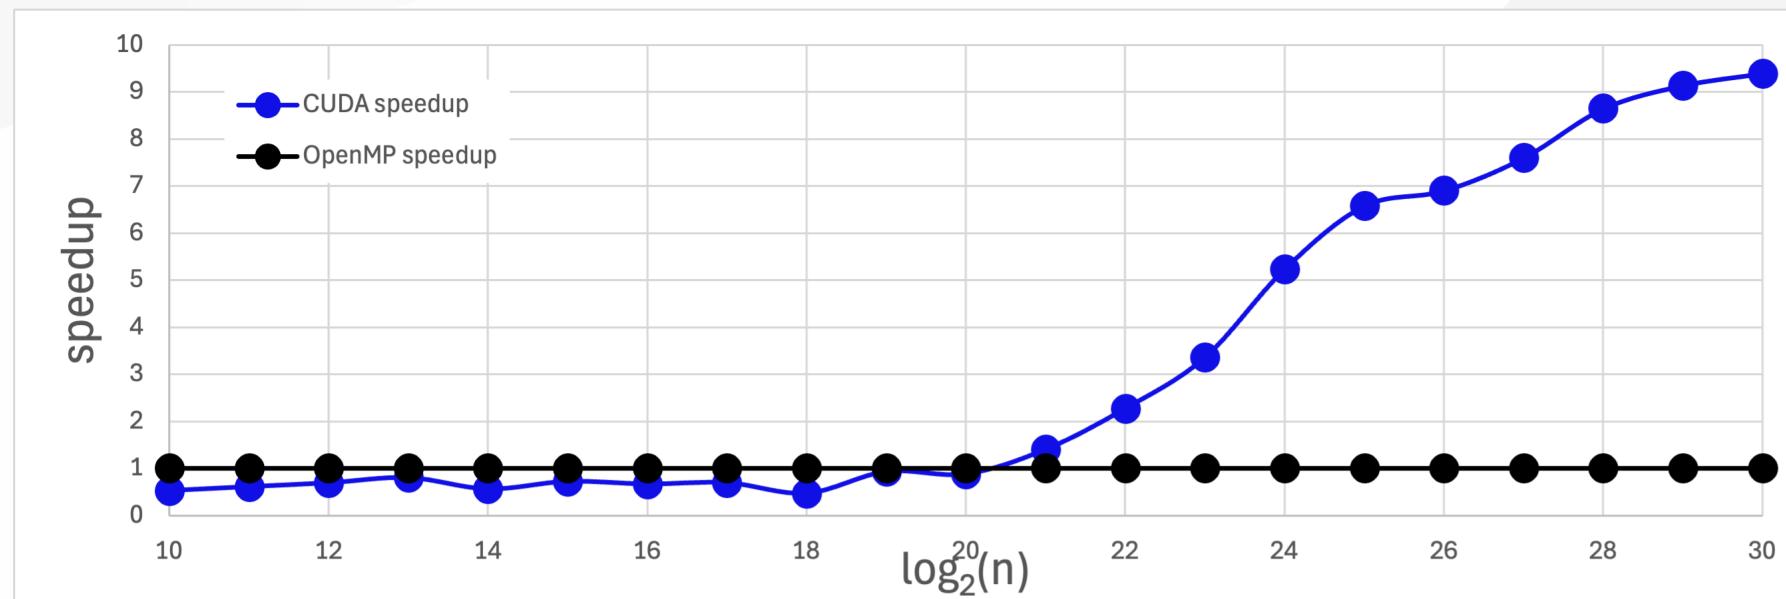
The GPU is a throughput device:

- CUDA breaks even for $n \geq 2^{21} \approx 2,000,000$
- requires $2^{29} \approx 540,000,000$ to gain "full" $9\times$ speedup

You have to provide enough parallelism to exploit many cores.

Exercise: Results

GH200



P100

